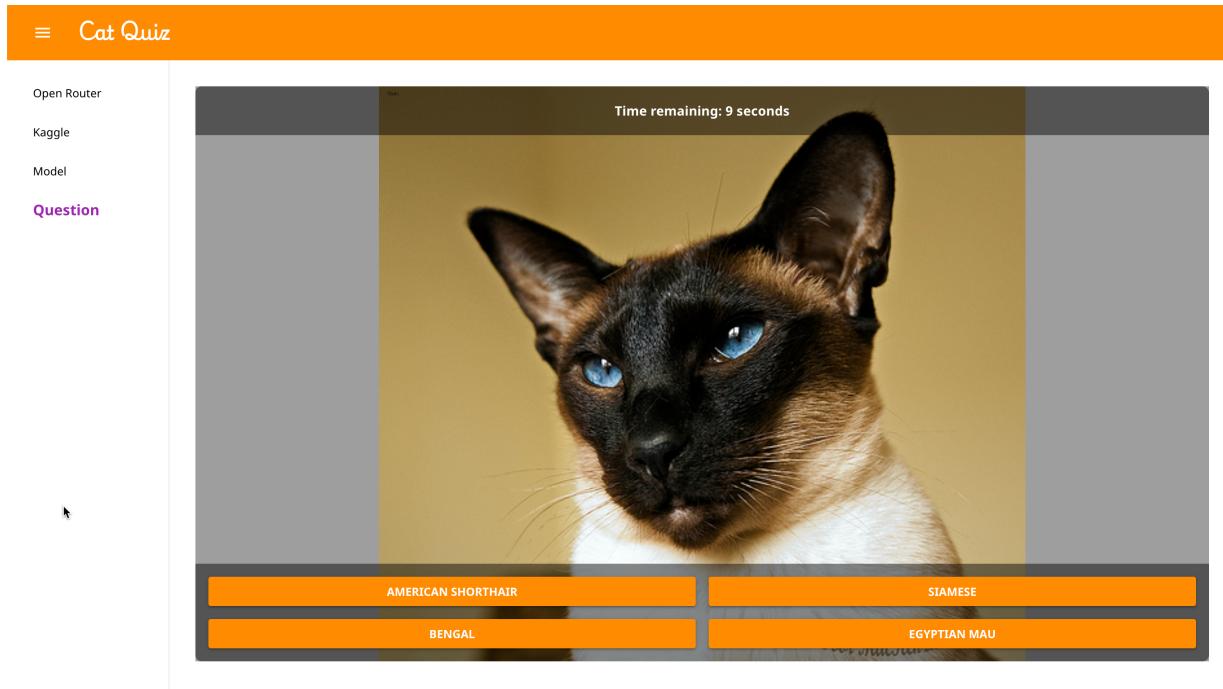
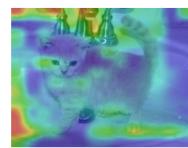
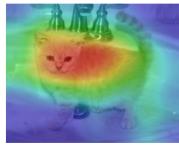
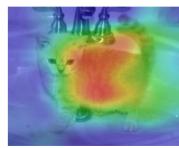
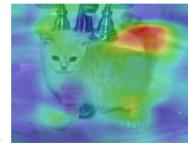


Cat Breed Classifier

This project is to train image classifier model to classify Cat Breeds dataset images. It has frontend as web application showing Quiz which allow user to take on guessing cat breed on the images.



Then after user pick answer, it will show result comparing user's with ChatGPT-4o-mini (OpenRouter API key required) and 4 Deep Learning Models. The rationale of LLM and GradCAM has been implemented to show model interpretability.

<p>Your Answer British Shorthair ✓ Correct!</p>	<p>ChatGPT-4o-mini Answer (1.53 seconds) British Shorthair ✓ Correct! The cat appears to have a stocky build and round face with short fur, which are characteristics typical of the British Shorthair breed.</p>	<p>Classification Model Answer (0.10 seconds) I don't know ✗ Incorrect! The correct answer was: British Shorthair </p>
<p>ResNet Model Answer (0.11 seconds) Persian ✗ Incorrect! The correct answer was: British Shorthair </p>	<p>EfficientNet Model Answer (0.43 seconds) British Shorthair ✓ Correct! </p>	<p>VGG Model Answer (0.13 seconds) I don't know ✗ Incorrect! The correct answer was: British Shorthair </p>

You : 2 ChatGPT: 2 Classification: 0 ResNet: 1 EfficientNet: 2 VGG: 1

> NEXT QUESTION

Data Sets

- Dataset is from Kaggle from ‘Geno Cat Breed Image Collection’ dataset (url: <https://www.kaggle.com/datasets/shawngano/gano-cat-breed-image-collection>)
- Contains 15 cat breeds with 375 photos for each breed (total 5,625 photos)
- Preprocessing step
 - Resize to 256x256
 - Random Crop to 224x224
 - Random Horizontal Flip
 - Random Rotation
 - Convert to Tensor
 - Normalize with mean and std of ImageNet

```
[1]: import os
import shutil
import kagglehub
from torchvision.datasets import ImageFolder
from torch.utils.data import DataLoader
import torchvision.models as models
from torchvision import transforms
import torch

def get_transformation() -> transforms.Compose:
    return transforms.Compose(
        [transforms.Resize((256, 256)),
         transforms.RandomCrop((224, 224)),
         transforms.RandomHorizontalFlip(),
         transforms.RandomRotation(10),
         transforms.ToTensor(),
         transforms.Normalize(mean=[0.485, 0.456, 0.406],
                             std=[0.229, 0.224, 0.225])]
    )

path = kagglehub.dataset_download("shawngano/gano-cat-breed-image-collection")

# Copy data to local folder using a relative path
local_path = "./data"

# Check if directory exists before removing it
if os.path.exists(local_path):
    shutil.rmtree(local_path)

# Copy the data (this will create the directory)
shutil.copytree(path, local_path)
path = local_path

transform = get_transformation()

dataset = ImageFolder(root=os.path.join(path, "Gano-Cat-Breeds-V1_1"), transform=transform)
loader = DataLoader(dataset, batch_size=32, shuffle=False, num_workers=8)
```

```
device = torch.device("mps" if torch.mps.is_available() else "cuda" if torch.cuda.  
is_available() else "cpu")
```

Model Architecture

Custom Model (1 Model)

1. CNN Model

- 5 Convolutional Layers
- 2 Fully Connected Layers
- 1 Dropout Layer in between the two FC Layers

from src/model_train.py

```
[2]: import os  
import time  
import torch  
import torch.nn as nn  
import torch.optim as optim  
from torch.utils.data import DataLoader, Dataset  
import torchvision  
import torchvision.transforms as transforms  
  
class CatBreedClassifier(nn.Module):  
    def __init__(self):  
        super(CatBreedClassifier, self).__init__()  
        self.conv_block = torch.nn.Sequential(  
            torch.nn.Conv2d(3, 64, kernel_size=3, stride=1, padding=1),  
            torch.nn.ReLU(),  
            torch.nn.MaxPool2d(kernel_size=2, stride=2), # Output: 64 x 112 x 112  
            torch.nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=1),  
            torch.nn.ReLU(),  
            torch.nn.MaxPool2d(kernel_size=2, stride=2), # Output: 128 x 56 x 56  
            torch.nn.Conv2d(128, 256, kernel_size=3, stride=1, padding=1),  
            torch.nn.ReLU(),  
            torch.nn.MaxPool2d(kernel_size=2, stride=2), # Output: 256 x 28 x 28  
            torch.nn.Conv2d(256, 512, kernel_size=3, stride=1, padding=1),  
            torch.nn.ReLU(),  
            torch.nn.MaxPool2d(kernel_size=2, stride=2), # Output: 512 x 14 x 14  
            torch.nn.Conv2d(512, 1024, kernel_size=3, stride=1, padding=1), # Changed  
            from 1024 to 512  
            torch.nn.ReLU(),  
            torch.nn.MaxPool2d(kernel_size=2, stride=2) # Output: 1024 x 7 x 7  
        )  
  
        # Calculate the flattened size: 512 channels * 7 * 7 = 25088  
        self.fc_block = torch.nn.Sequential(  
            torch.nn.Flatten(),  
            torch.nn.Linear(1024 * 7 * 7, 512), # Adjusted input size  
            torch.nn.ReLU(),  
            torch.nn.Dropout(0.5),  
            torch.nn.Linear(512, 15), # 15 breeds
```

```
)
```

```
def forward(self, x):
    x = self.conv_block(x)
    x = self.fc_block(x)
    return x
```

Transfer Learning (3 models)

2. VGG 16 Model

```
[3]: def get_vgg_model() -> nn.Module:
    vgg_model = torchvision.models.vgg16(weights=torchvision.models.VGG16_Weights.
    ↪IMAGENET1K_V1)
    for param in vgg_model.parameters():
        param.requires_grad = False
    vgg_model.classifier[6] = nn.Linear(vgg_model.classifier[6].in_features, 15)
    return vgg_model
```

3. ResNet 18 Model

```
[4]: def get_resnet_model() -> nn.Module:
    res_model = torchvision.models.resnet18(weights=torchvision.models.
    ↪ResNet18_Weights.IMAGENET1K_V1)
    for param in res_model.parameters():
        param.requires_grad = False
    res_model.fc = nn.Linear(res_model.fc.in_features, 15)
    return res_model
```

4. EfficientNet B2 Model

```
[5]: def get_efficient_net_model() -> nn.Module:
    eff_model = torchvision.models.efficientnet_b2(weights=torchvision.models.
    ↪EfficientNet_B2_Weights.IMAGENET1K_V1)
    for param in eff_model.parameters():
        param.requires_grad = False
    eff_model.classifier[1] = nn.Linear(eff_model.classifier[1].in_features, 15)
    return eff_model
```

Training

All 4 models has been trained using same hyperparamters for comparison.

- Training Batch Size: 128
- Epochs: 20 for CNN model and 10 for transfer learning models.
- Loss Function: Cross Entropy Loss
- Optimizer: SGD
 - Learning Rate: 0.01 for CNN model and 0.001 for transfer learning models.
 - Momentum: 0.9

After finish training model saved to .pth for later used for inference result on Cat Quiz program.

classification_model.pth	127.9 MB	Today at 8:41AM	Document
efficientnet_model.pth	31.6 MB	Today at 10:06AM	Document
resnet_model.pth	44.9 MB	Today at 10:52AM	Document
vgg_model.pth	537.3 MB	Today at 9:03AM	Document

- train_model = function to train model and saved training accuracy/loss and validation accuracy/loss for later display.

```
[6]: def train_model(app_state, model: nn.Module, model_name: str) -> tuple[list[float], list[float], list[float], float]:
    print(f"Training model on {app_state.device}")
    transform = get_transformation()

    print(f"Loading dataset")
    dataset = ImageFolder(app_state.image_path, transform=transform)

    # Split dataset into train and test sets (80% train, 20% test)
    train_size = int(0.8 * len(dataset))
    test_size = len(dataset) - train_size
    train_dataset, test_dataset = torch.utils.data.random_split(dataset, [train_size, test_size])

    # Create data loaders for both sets
    train_loader = DataLoader(train_dataset, batch_size=128, shuffle=True, num_workers=4)
    test_loader = DataLoader(test_dataset, batch_size=128, shuffle=False, num_workers=4)

    print(f"Dataset split: {train_size} training samples, {test_size} test samples")

    print(f"Loading model")
    model.to(app_state.device)
    criterion = nn.CrossEntropyLoss()
    optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)

    training_loss = []
    training_accuracy = []
    validation_loss = []
    validation_accuracy = []

    start_time = time.time()
    for epoch in range(5):
        model.train()

        running_loss = 0.0
        correct = 0
        total = 0
        for i, data in enumerate(train_loader, 0):
            inputs, labels = data
            inputs = inputs.to(app_state.device)
```

```

        labels = labels.to(app_state.device)
        optimizer.zero_grad()

        outputs = model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()
        _, predicted = torch.max(outputs.data, 1)
        correct += (predicted == labels).sum().item()
        total += labels.size(0)

        print(f"Epoch {epoch + 1}, Batch {i + 1}, Loss: {loss.item()}")

    training_loss.append(running_loss / len(train_loader))
    training_accuracy.append(correct / total)

    model.eval()
    with torch.no_grad():
        running_loss = 0.0
        correct = 0
        total = 0
        for i, data in enumerate(test_loader, 0):
            inputs, labels = data
            inputs = inputs.to(app_state.device)
            labels = labels.to(app_state.device)
            outputs = model(inputs)
            loss = criterion(outputs, labels)
            running_loss += loss.item()
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

        validation_loss.append(running_loss / len(test_loader))
        validation_accuracy.append(correct / total)

        print(f"Accuracy of the network on the {total} test images: {100 * correct / total}%")
    print(f"Finished Training")
    print(f"Loss: {running_loss / len(train_loader)}")
    training_time = time.time() - start_time
    print(f"Saving model to {model_name}")
    scripted_model = torch.jit.script(model)
    torch.jit.save(scripted_model, model_name)


    return training_loss, training_accuracy, validation_loss, validation_accuracy, training_time

```

Result and Evaluation


Training Accuracy/Loss

- While models were trained directly in application, both training and validation accuracy/loss will be kepted for visualization in web application later. Below chart shows training and validation accuracy/loss for each model that captured directly from web application.



- Chart below shows time used for training each model, custom CNN train for 20 epochs, while others model trained for 10 epochs.

Training Time (seconds)



Testing Accuracy/Loss

```
[7]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, classification_report
import matplotlib.pyplot as plt

def predict_and_evaluate(model):
    # For collecting all targets and predictions
    all_preds = []
    all_labels = []

    model.eval()

    with torch.no_grad():
        for images, labels in loader:
            images = images.to(device)
            labels = labels.to(device)

            outputs = model(images)
            _, preds = torch.max(outputs, 1)

            all_preds.extend(preds.cpu().numpy())
            all_labels.extend(labels.cpu().numpy())

    plt.figure(figsize=(10, 7))
    cm = confusion_matrix(all_labels, all_preds)
```

```

disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=dataset.classes)
disp.plot()

plt.title(f"Confusion Matrix - {model.__class__.__name__}")
plt.xticks(rotation=90)
plt.tight_layout()
plt.show()

print(classification_report(all_labels, all_preds, target_names=dataset.classes))


```

CNN Model

```
[8]: cnn_model = CatBreedClassifier().to(device)
cnn_model.load_state_dict(torch.jit.load('classification_model.pth').state_dict())

predict_and_evaluate(cnn_model)
```

<Figure size 1000x700 with 0 Axes>



	precision	recall	f1-score	support
Abyssinian	0.35	0.25	0.29	375
American Bobtail	0.24	0.06	0.09	375
American Shorthair	0.28	0.10	0.15	375
Bengal	0.54	0.54	0.54	375


Birman	0.43	0.46	0.45	375
Bombay	0.66	0.67	0.66	375
British Shorthair	0.18	0.14	0.16	375
Egyptian Mau	0.46	0.47	0.46	375
Maine Coon	0.44	0.18	0.26	375
Persian	0.45	0.29	0.35	375
Ragdoll	0.37	0.60	0.46	375
Russian Blue	0.31	0.77	0.44	375
Siamese	0.51	0.47	0.49	375
Sphynx	0.32	0.45	0.37	375
Tuxedo	0.66	0.86	0.74	375
accuracy			0.42	5625
macro avg	0.41	0.42	0.39	5625
weighted avg	0.41	0.42	0.39	5625

EfficientNet B2 Model

```
[9]: eff_model = get_efficient_net_model().to(device)
eff_model.load_state_dict(torch.jit.load('efficientnet_model.pth').state_dict())

predict_and_evaluate(eff_model)
```

<Figure size 1000x700 with 0 Axes>




	precision	recall	f1-score	support
Abyssinian	0.71	0.70	0.71	375
American Bobtail	0.47	0.18	0.26	375
American Shorthair	0.52	0.38	0.44	375
Bengal	0.63	0.78	0.70	375
Birman	0.68	0.76	0.72	375
Bombay	0.76	0.93	0.84	375
British Shorthair	0.72	0.40	0.52	375
Egyptian Mau	0.65	0.70	0.67	375
Maine Coon	0.69	0.71	0.70	375
Persian	0.77	0.91	0.83	375
Ragdoll	0.70	0.54	0.61	375
Russian Blue	0.71	0.86	0.78	375
Siamese	0.83	0.91	0.87	375
Sphynx	0.83	0.97	0.89	375
Tuxedo	0.78	0.97	0.86	375
accuracy			0.71	5625
macro avg	0.70	0.71	0.69	5625
weighted avg	0.70	0.71	0.69	5625

ResNet18 Model

```
[10]: res_model = get_resnet_model().to(device)
res_model.load_state_dict(torch.jit.load('resnet_model.pth').state_dict())
predict_and_evaluate(res_model)
```

<Figure size 1000x700 with 0 Axes>

Confusion Matrix - ResNet




	precision	recall	f1-score	support
Abyssinian	0.81	0.67	0.73	375
American Bobtail	0.45	0.33	0.38	375
American Shorthair	0.50	0.47	0.48	375
Bengal	0.76	0.76	0.76	375
Birman	0.70	0.75	0.72	375
Bombay	0.79	0.92	0.85	375
British Shorthair	0.71	0.59	0.64	375
Egyptian Mau	0.67	0.74	0.70	375
Maine Coon	0.70	0.71	0.71	375
Persian	0.84	0.90	0.87	375
Ragdoll	0.66	0.63	0.64	375
Russian Blue	0.78	0.81	0.79	375
Siamese	0.87	0.84	0.85	375
Sphynx	0.84	0.96	0.90	375
Tuxedo	0.82	0.95	0.88	375
accuracy			0.73	5625
macro avg	0.73	0.73	0.73	5625
weighted avg	0.73	0.73	0.73	5625

VGG Model

```
[11]: vgg_model = get_vgg_model().to(device)
vgg_model.load_state_dict(torch.jit.load('vgg_model.pth').state_dict())

predict_and_evaluate(vgg_model)
```

<Figure size 1000x700 with 0 Axes>



	precision	recall	f1-score	support
Abyssinian	0.75	0.76	0.75	375
American Bobtail	0.53	0.31	0.39	375
American Shorthair	0.48	0.52	0.50	375
Bengal	0.74	0.78	0.76	375
Birman	0.74	0.71	0.72	375
Bombay	0.82	0.95	0.88	375
British Shorthair	0.75	0.57	0.65	375
Egyptian Mau	0.70	0.75	0.72	375
Maine Coon	0.79	0.81	0.80	375
Persian	0.91	0.89	0.90	375
Ragdoll	0.65	0.75	0.69	375
Russian Blue	0.79	0.82	0.81	375
Siamese	0.87	0.89	0.88	375
Sphynx	0.95	0.94	0.94	375

Tuxedo	0.85	0.94	0.89	375
accuracy			0.76	5625
macro avg	0.75	0.76	0.75	5625
weighted avg	0.75	0.76	0.75	5625




Visualization and Interpretability

- GradCAM has been implemented for each models.
 - Color will gradient from Blue(0%) -> Cyan(25%) -> Green(50%) -> Yellow(75%) -> Red(100%) impact to features that CNN layer consider.
- Code below show GradCAM implemented on CNN Layer selection for each model from src\model_train.py

```
target_layer = None
```

```
if model_name == Models.ClassificationModel:
    # For custom CNN, use the last conv layer in conv_block
    target_layer = original_model.conv_block[-3] # Get the last Conv2d before the final MaxPool
elif model_name == Models.ResNetModel:
    # For ResNet, use the last layer in layer4
    target_layer = original_model.layer4[-1].conv2
elif model_name == Models.EfficientNetModel:
    # For EfficientNet, use the last conv layer in features
    target_layer = original_model.features[-1][0]
elif model_name == Models.VGGModel:
    # For VGG, use the last conv layer in features
    for module in original_model.features:
        if isinstance(module, nn.Conv2d):
            target_layer = module
```

Below image show GradCAM implemented in web application showing each model's area of interest on a cat image.

Your Answer Abyssinian ✕ Incorrect! The correct answer was: Birman	ChatGPT-4o-mini Answer (2.33 seconds) Birman ✓ Correct! The cat has a characteristic color-point pattern and blue eyes, which are typical features of the Birman breed.	Classification Model Answer (0.50 seconds) Birman ✓ Correct! 
ResNet Model Answer (0.11 seconds) Birman ✓ Correct! 	EfficientNet Model Answer (0.44 seconds) Birman ✓ Correct! 	VGG Model Answer (0.27 seconds) Birman ✓ Correct! 

Conclusion

- In this project, we have created a web application using NiceGUI framework. This application uses PyTorch to train and evaluate models for different cat breeds. Custom CNN, ResNet18, EfficientNet B2, and VGG16 models.
- GradCAM has been implemented for each models to visualize the feature importance of the models.

Model	Parameters	Accuracy	Loss	Training Time (seconds)
Custom CNN Model	31.9 Millions	42%	1.86	2,698
EfficientNet Model	9.1 Millions	71%	1.40	1,256
ResNet Model	11.7 Millions	73%	0.99	2,357
VGG Model	138.3 Millions	76%	0.81	1,314

- VGG16 Model is the best model bases on confusion matrix and accuracy score, but it contains very large number of parameters (138 Million parameters).
- ResNet18 Model on the other hand, has smaller number of parameters (11.7 Million parameters) and nearly the same performance as VGG16.
- EfficientNet B2 has 9.1 Million parameters and slightly lower accuracy compared to VGG16 and ResNet18 on this dataset.
- ChatGTP-4o-mini which is multi-model has image recognition capabilities. It can be used for cat image classification and provide rationale for its predictions.